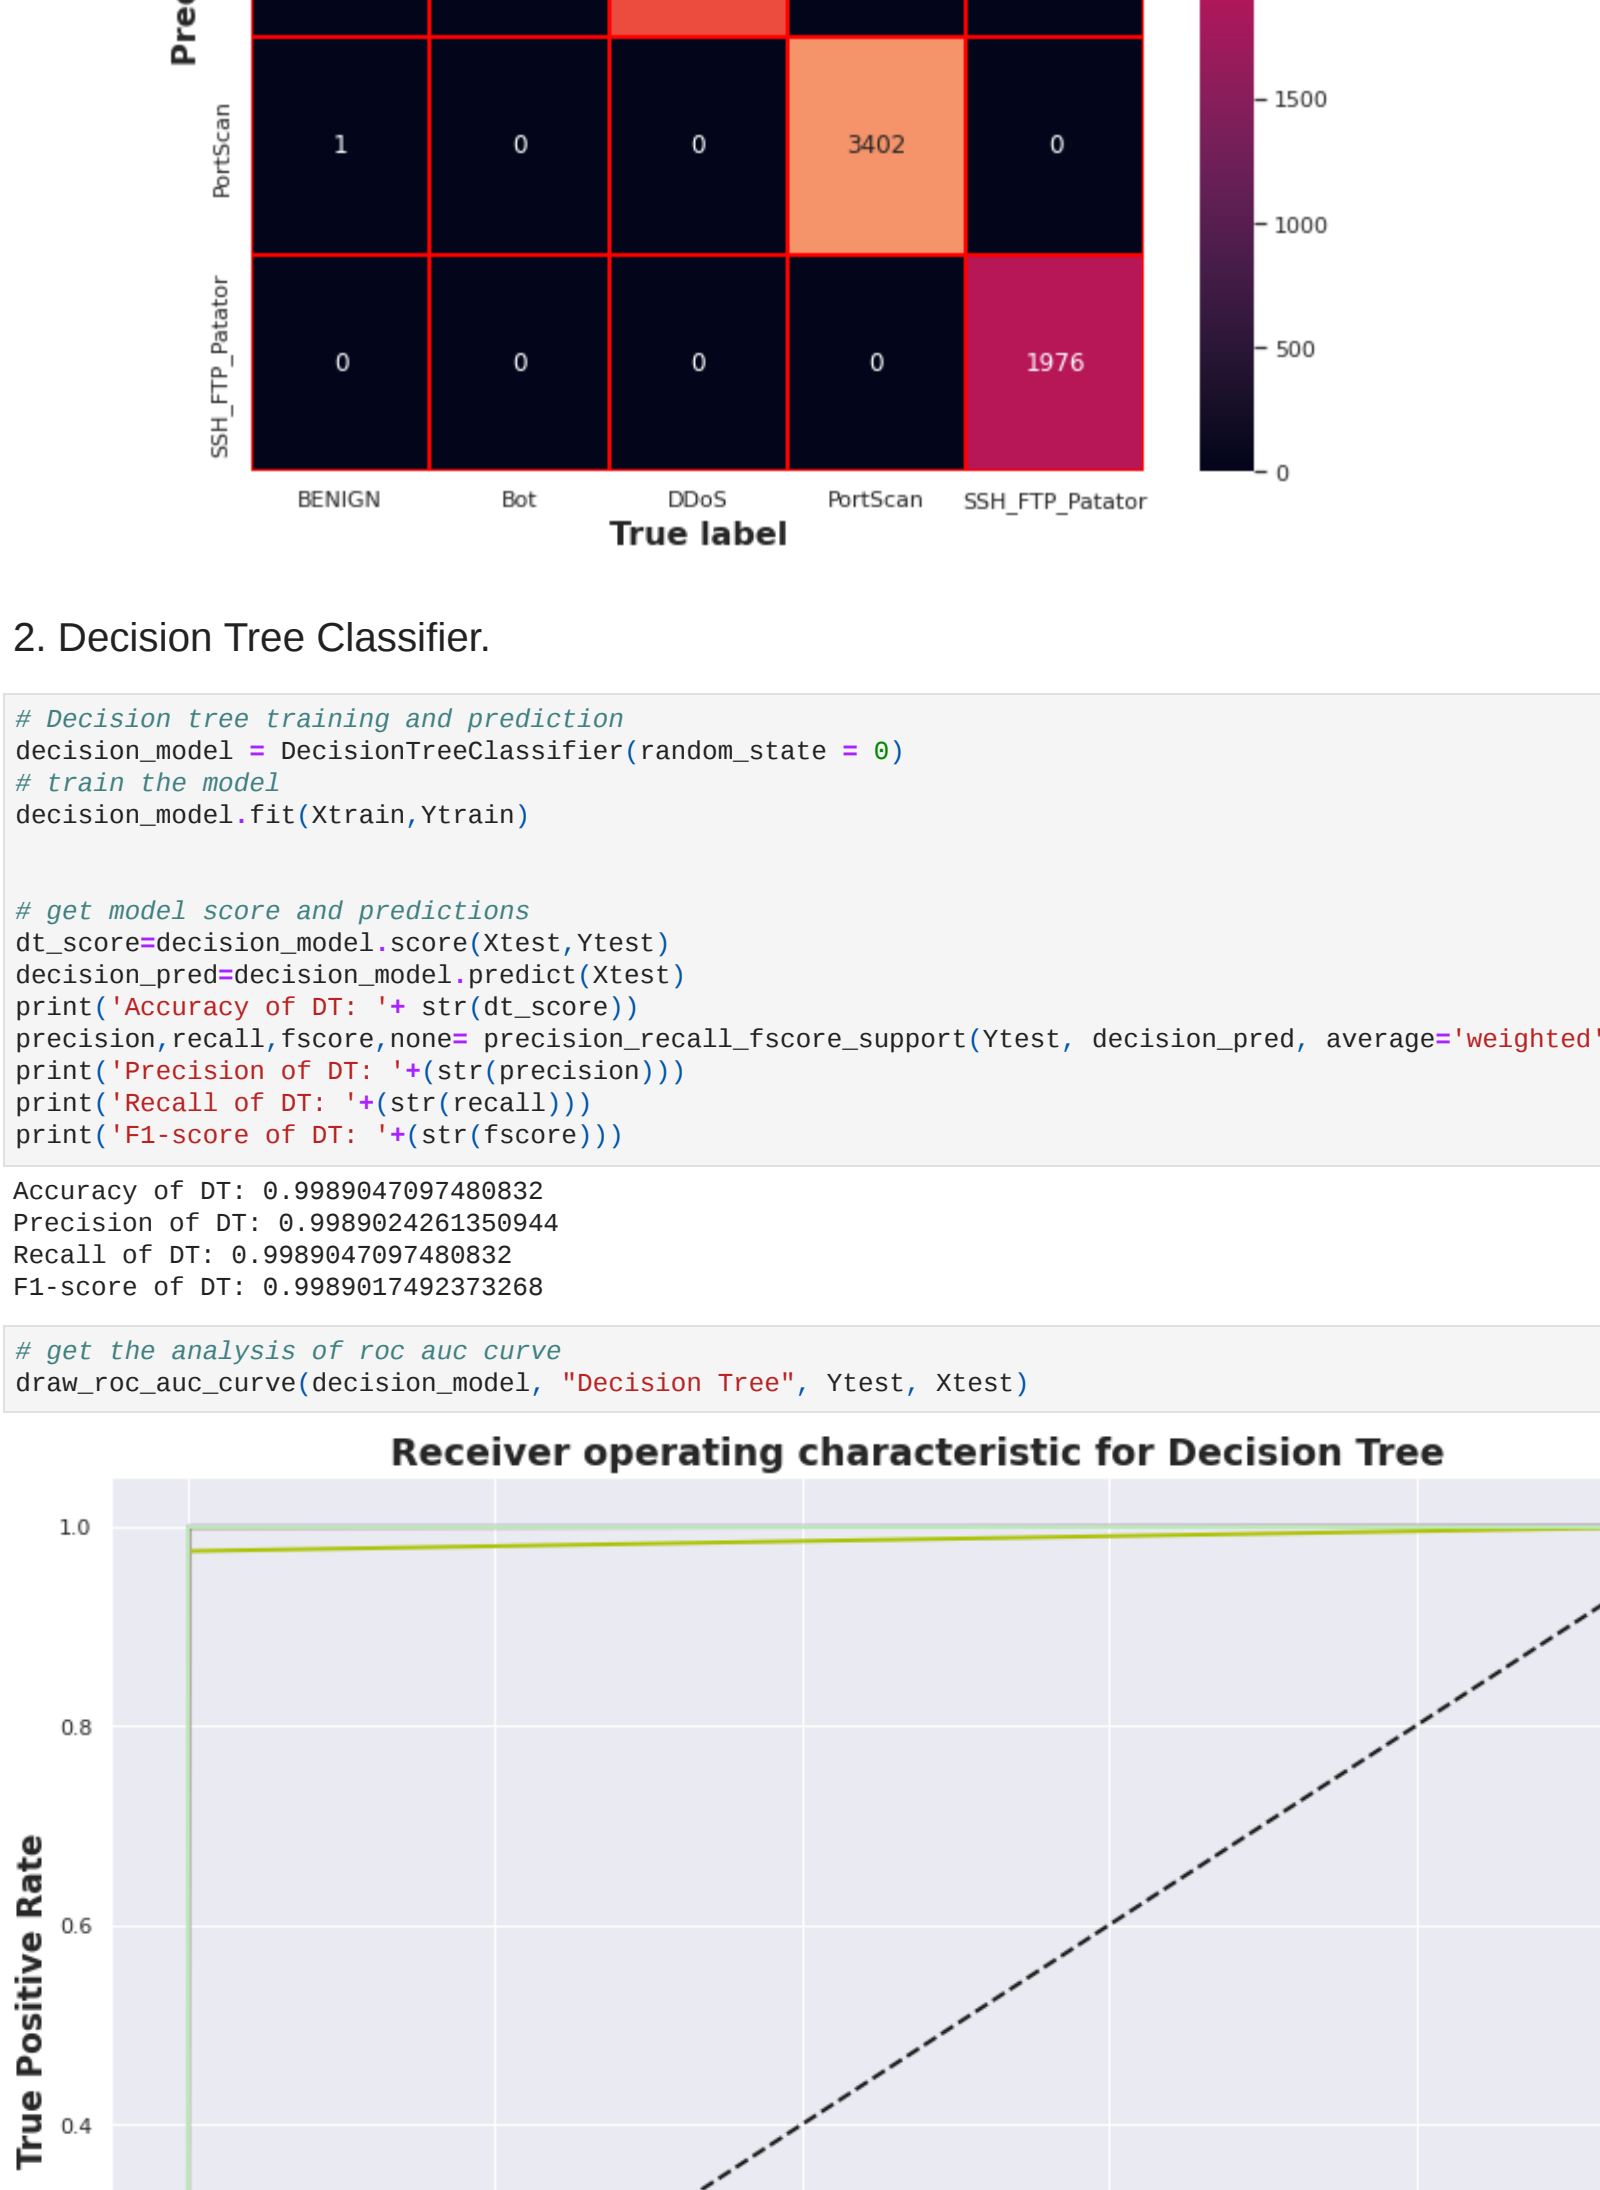


Accuracy is 0.9989047097480832
F1 Score is 0.9989017492373268
Recall score is 0.9989047097480832
Precision Score is 0.9989024261350944

Outputting Classification Report Matrix Plot

	precision	recall	f1-score	support
BENIGN	1.00	1.00	1.00	4380
Bot	0.99	0.97	0.98	280
DDoS	1.00	1.00	1.00	2743
PortScan	1.00	1.00	1.00	3403
SSH_FTP_Patator	1.00	1.00	1.00	1976
accuracy		1.00	1.00	12782
macro avg	1.00	0.99	1.00	12782
weighted avg	1.00	1.00	1.00	12782

Confusion Plot for Logistic Regression Classifier's Performance



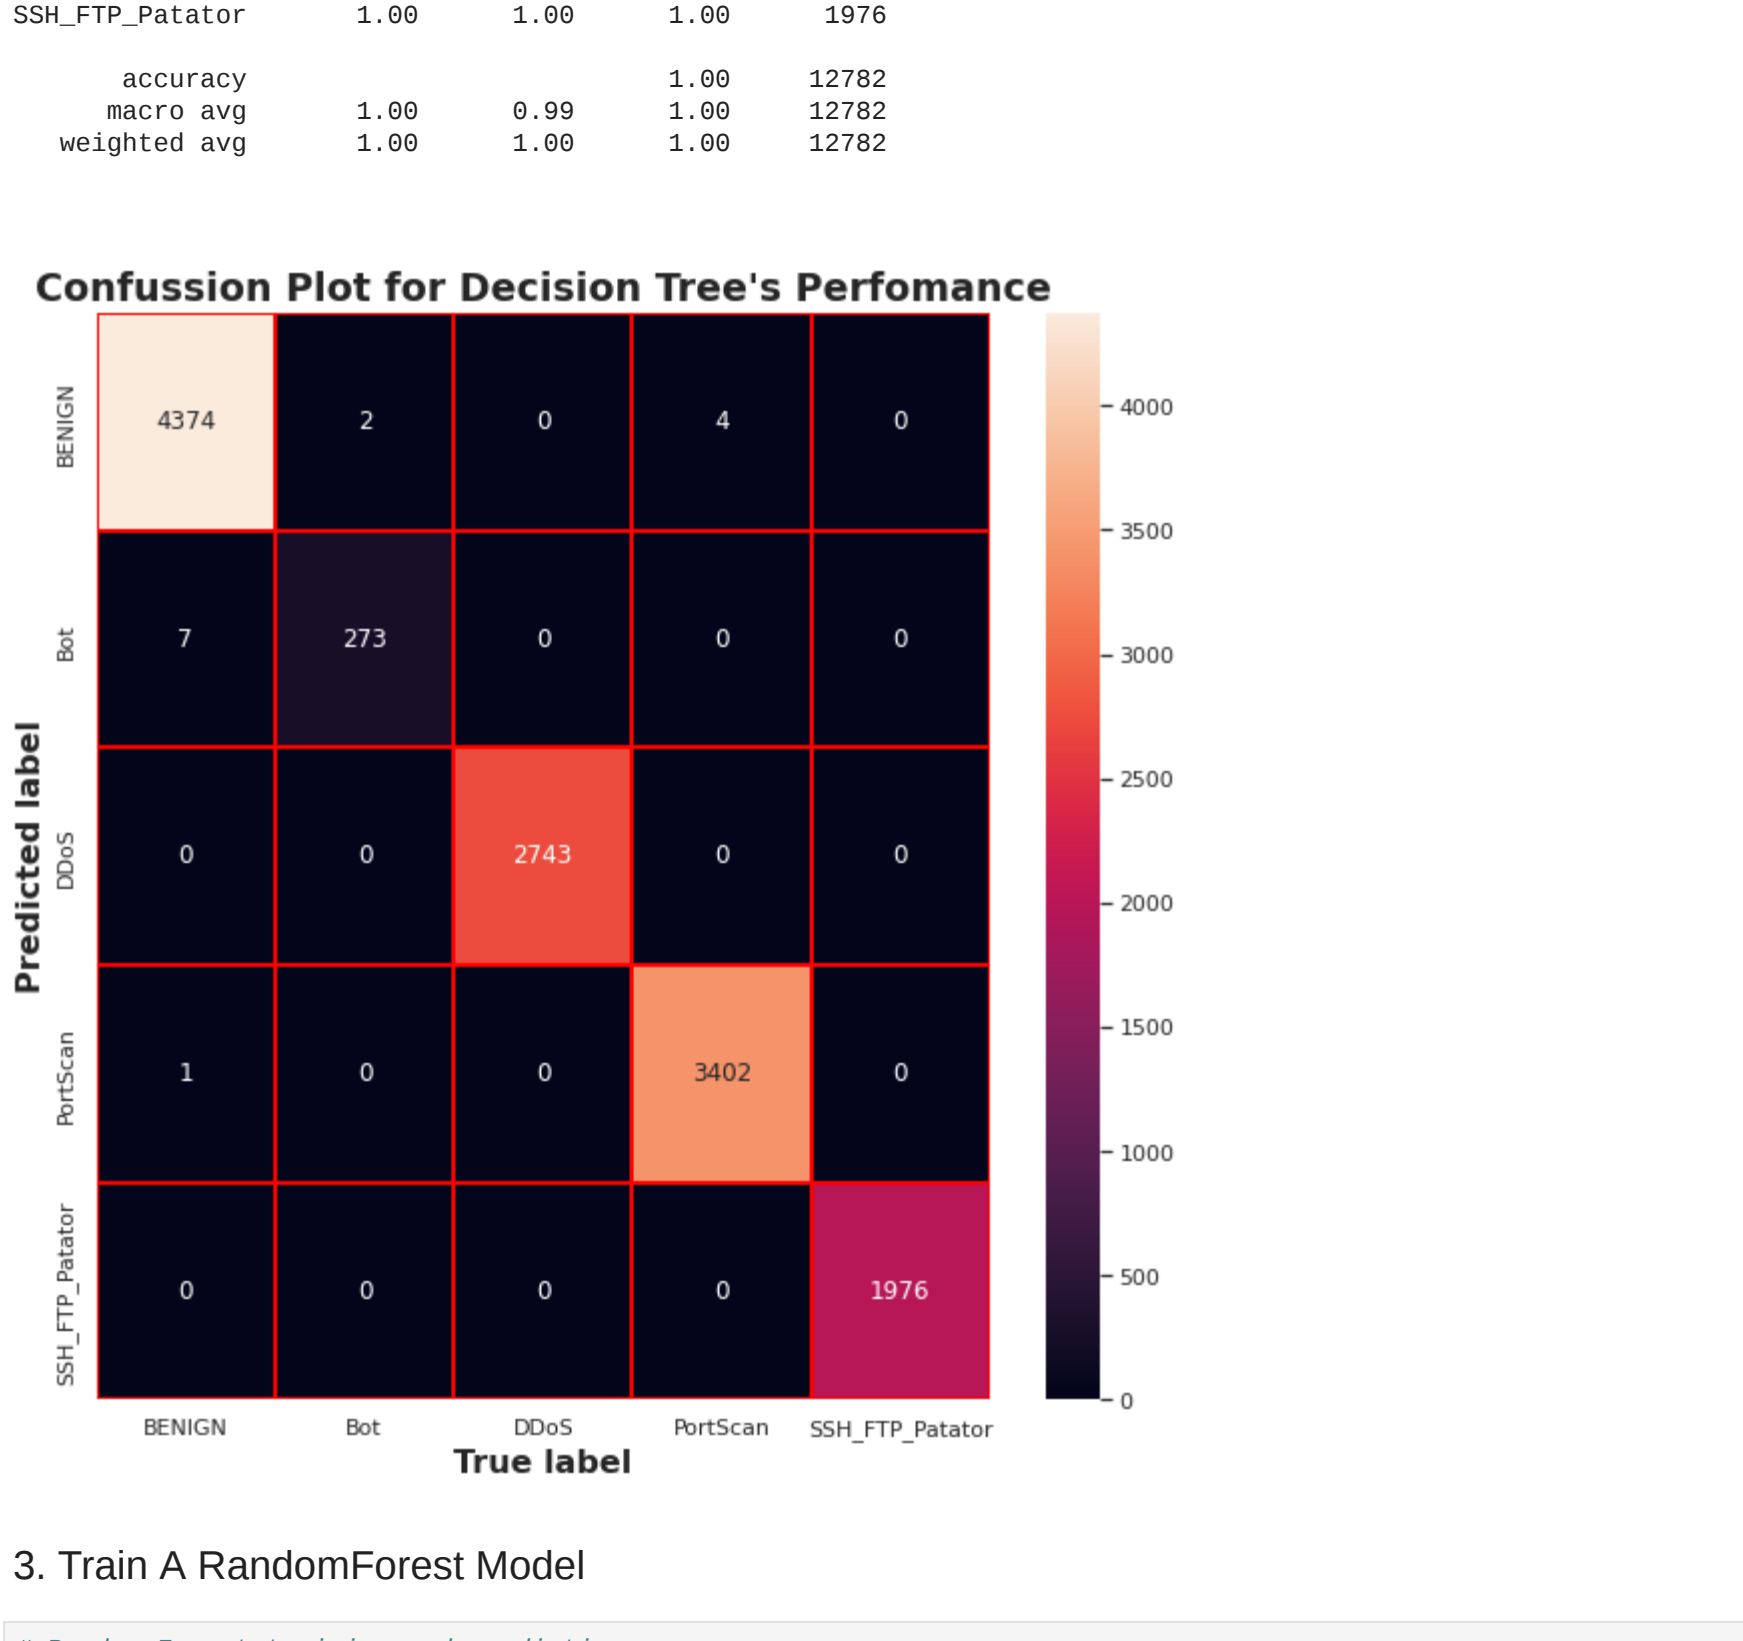
2. Decision Tree Classifier.

```
In [54]: # Decision tree training and prediction
decision_model = DecisionTreeClassifier(random_state = 0)
# train the model
decision_model.fit(X_train,Y_train)

# get model score and predictions
dt_score=decision_model.score(X_test,Ytest)
decision_pred=decision_model.predict(Xtest)
print('Accuracy of DT: %s'%(str(dt_score)))
precision,recall,f1score,none: precision_recall_fscore_support(Ytest, decision_pred, average='weighted')
print('Precision of DT: %s'%(str(precision)))
print('Recall of RF: %s'%(str(recall)))
print('F1-score of DT: %s'%(str(f1score)))

Accuracy of DT: 0.9989047097480832
Precision of DT: 0.9989024261350944
Recall of DT: 0.9989017492373268
F1-score of DT: 0.9989037492373268
```

```
In [55]: # get the analysis of roc auc curve
draw_roc_auc_curve(decision_model, "Decision Tree", Ytest, Xtest)
```



```
In [56]: # confusion matrix analysis
draw_confusion_matrix(Ytest, decision_model.predict(Xtest), "Decision Tree")
```

Decision Tree Metrics Analysis Results

Accuracy is 0.9989047097480832
F1 Score is 0.9989017492373268
Recall score is 0.9989047097480832
Precision Score is 0.9989024261350944

Outputting Classification Report Matrix Plot

	precision	recall	f1-score	support
BENIGN	1.00	1.00	1.00	4380
Bot	0.99	0.97	0.98	280
DDoS	1.00	1.00	1.00	2743
PortScan	1.00	1.00	1.00	3403
SSH_FTP_Patator	1.00	1.00	1.00	1976
accuracy		1.00	1.00	12782
macro avg	1.00	0.99	1.00	12782
weighted avg	1.00	1.00	1.00	12782

3. Train A RandomForest Model

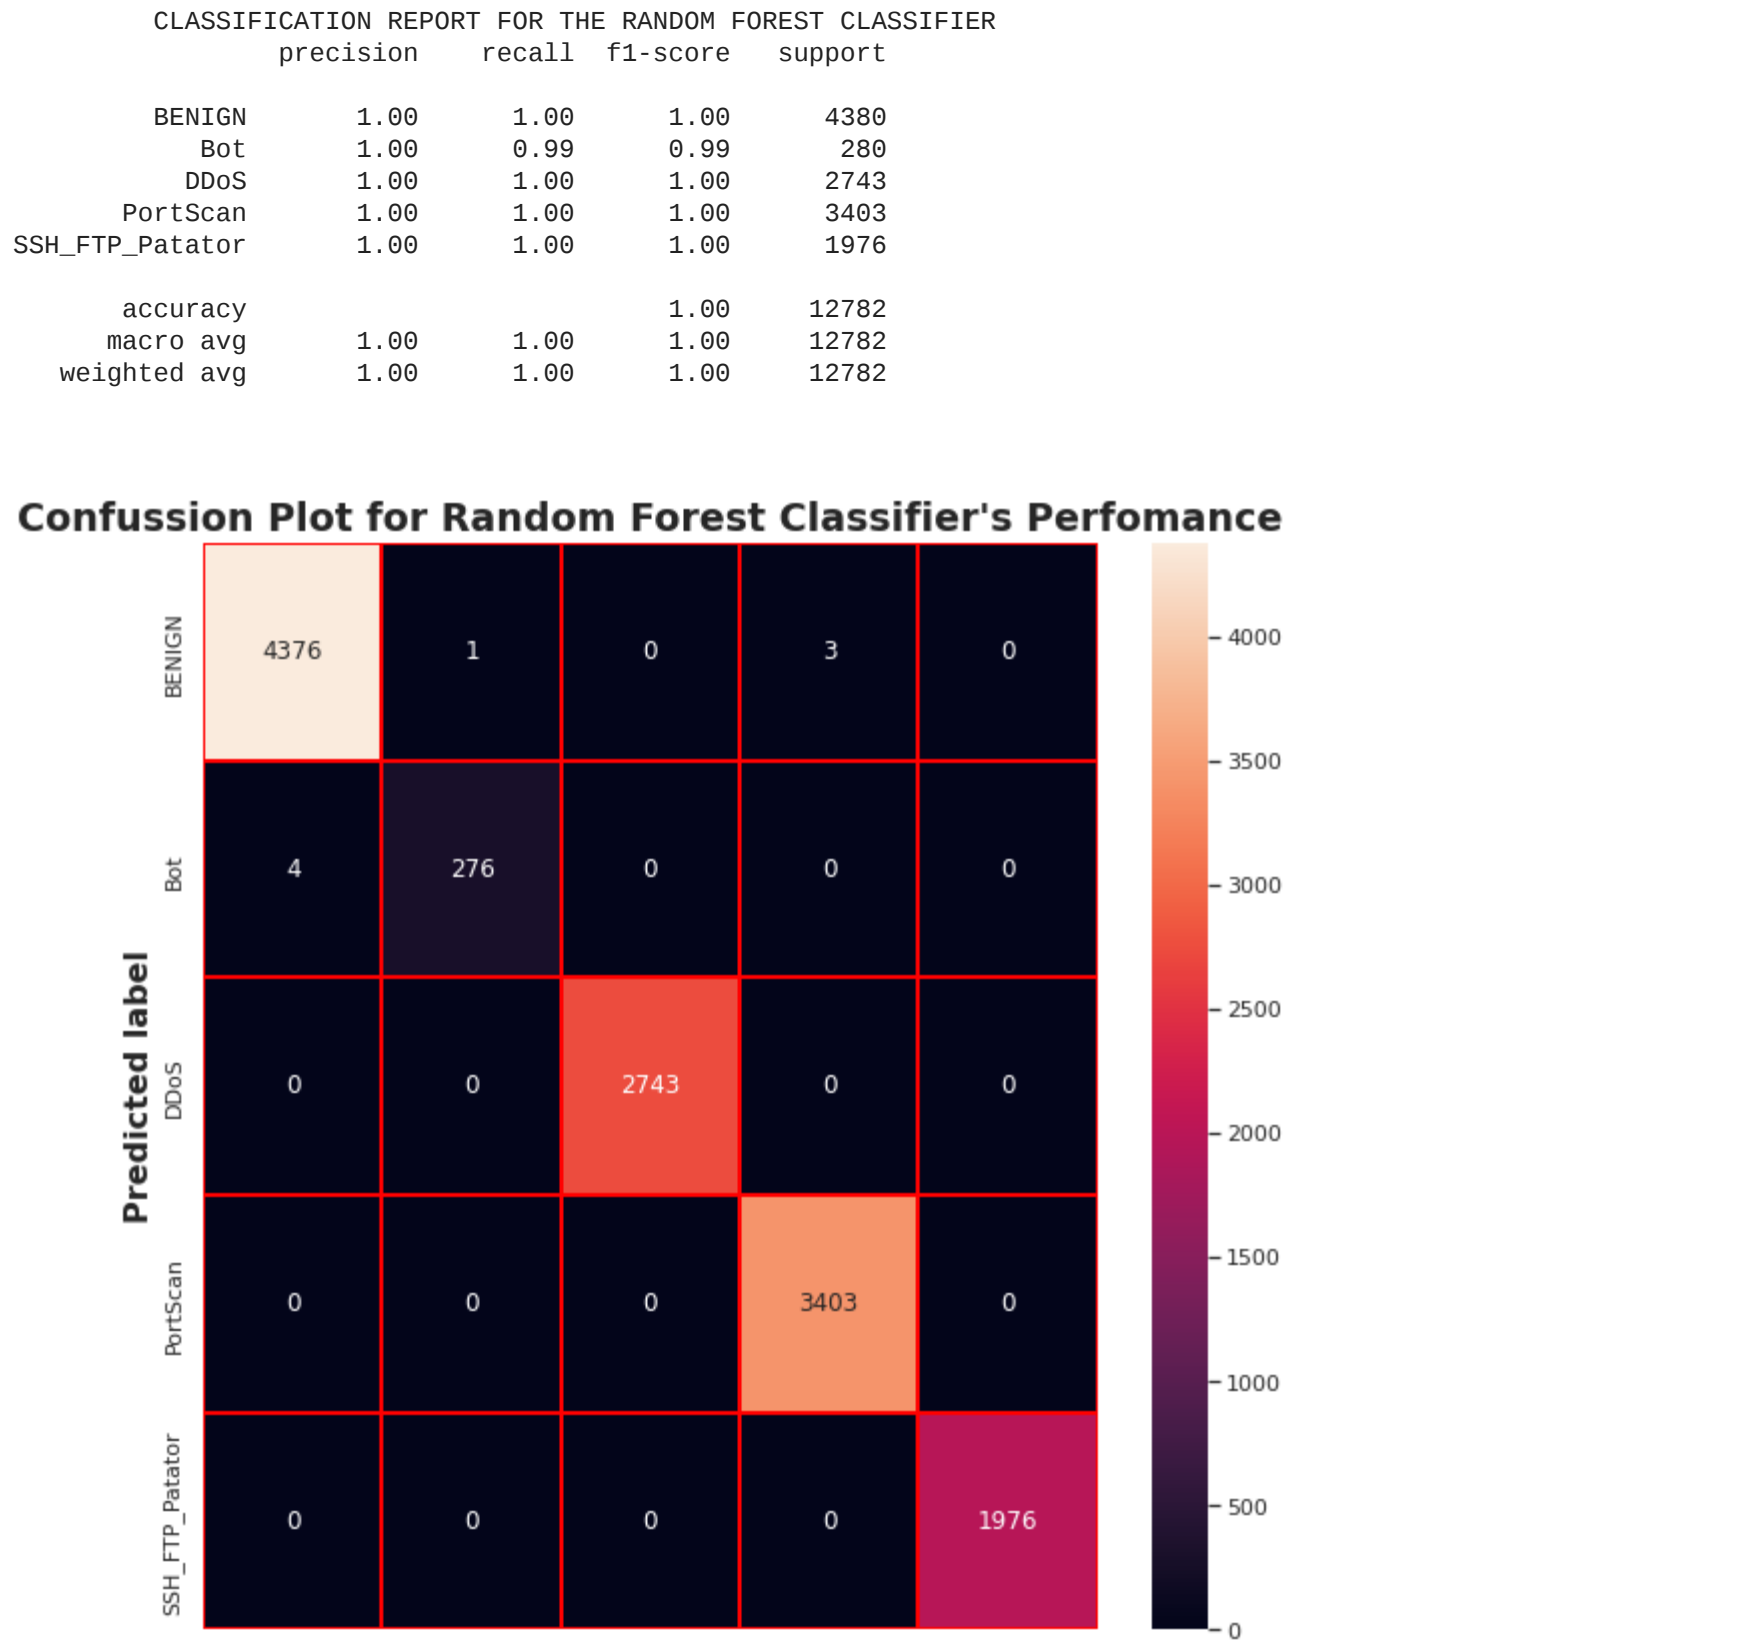
```
In [57]: # Random Forest training and prediction
rf_clf = RandomForestClassifier(random_state = 0)
rf_clf.fit(X_train,Y_train)
rf_score=rf_clf.score(Xtest,Ytest)
rf_clf_predict=rf_clf.predict(Xtest)
xgb_predict=xgb_clf.predict(Xtest)
print('Accuracy of RF: %s'%(str(rf_score)))
precision,recall,f1score,none: precision_recall_fscore_support(Ytest, rf_clf_predict, average='weighted')
print('Precision of XGBoost: %s'%(str(precision)))
print('Recall of RF: %s'%(str(recall)))
print('F1-score of RF: %s'%(str(f1score)))

Accuracy of RF: 0.9993741198560475
Precision of RF: 0.9993734792732641
Recall of RF: 0.9993741198560475
F1-score of RF: 0.9993731181269735
```

```
In [58]: print("Hey there")
```

Hey there

```
In [59]: # get the analysis of roc auc curve for random forest
draw_roc_auc_curve(rf_clf, "Random Forest Classifier", Ytest, Xtest)
```



```
In [60]: # confusion matrix analysis for random forest
draw_confusion_matrix(Ytest, rf_clf.predict(Xtest), "Random Forest Classifier")
```

Random Forest Classifier Metrics Analysis Results

Accuracy is 0.9993741198560475
F1 Score is 0.9993731181269735
Recall score is 0.9993741198560475
Precision Score is 0.9993734792732641

Outputting Classification Report Matrix Plot

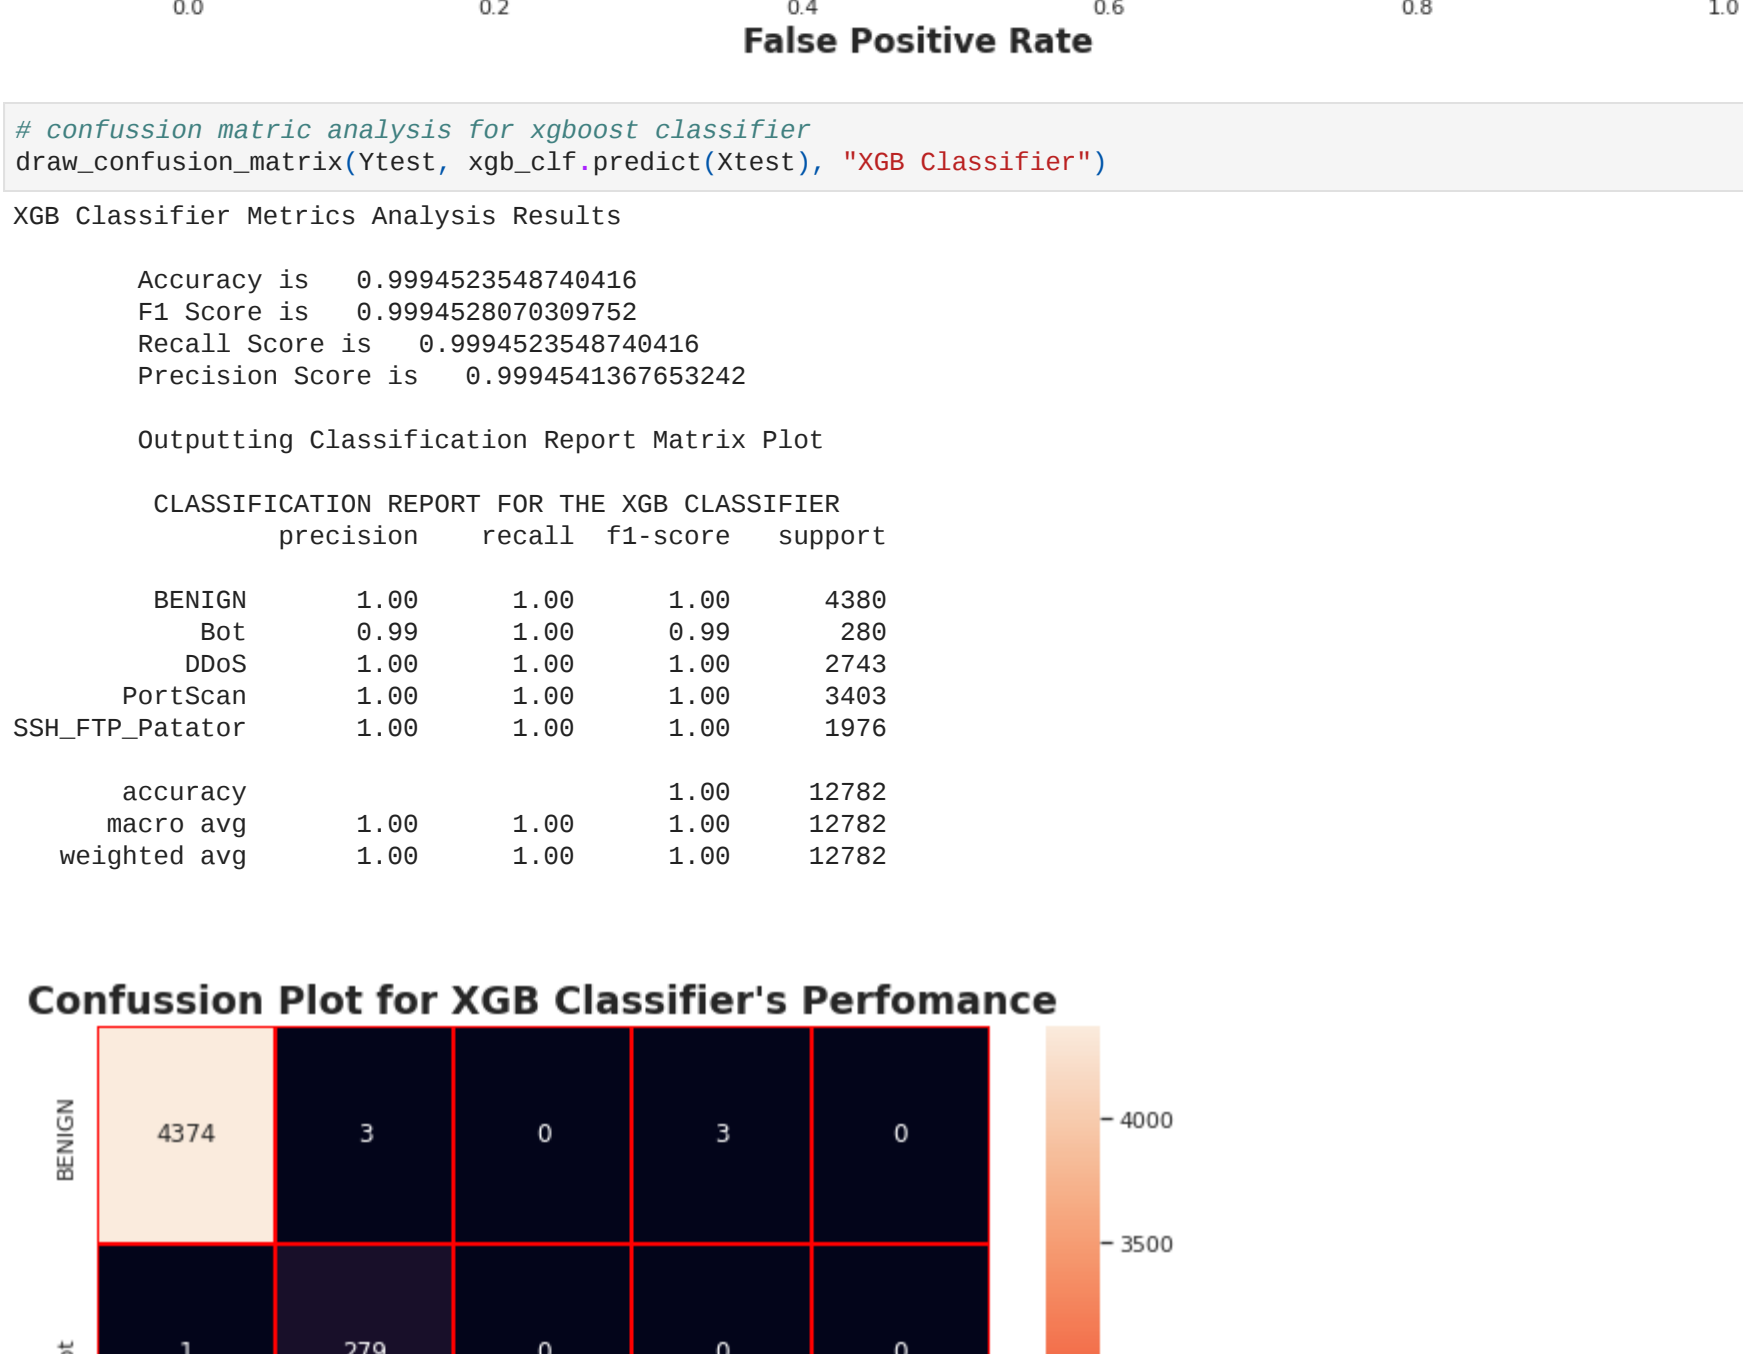
	precision	recall	f1-score	support
BENIGN	1.00	1.00	1.00	4380
Bot	1.00	0.99	0.99	280
DDoS	1.00	1.00	1.00	2743
PortScan	1.00	1.00	1.00	3403
SSH_FTP_Patator	1.00	1.00	1.00	1976
accuracy		1.00	1.00	12782
macro avg	1.00	1.00	1.00	12782
weighted avg	1.00	1.00	1.00	12782

4. XGBOOST Model

```
In [61]: # xgboost training and prediction
xgb_clf = XGBClassifier(n_estimators = 500)
xgb_clf.fit(X_train,Y_train)
xgb_score=xgb_clf.score(Xtest,Ytest)
xgb_predict=xgb_clf.predict(Xtest)
print('Accuracy of XGBoost: %s'%(str(xgb_score)))
precision,recall,f1score,none: precision_recall_fscore_support(Ytest, xgb_predict, average='weighted')
print('Precision of XGBoost: %s'%(str(precision)))
print('Recall of XGBoost: %s'%(str(recall)))
print('F1-score of XGBoost: %s'%(str(f1score)))

Accuracy of XGBoost: 0.9994523548740416
Precision of XGBoost: 0.999451367653242
Recall of XGBoost: 0.9994523548740416
F1-score of XGBoost: 0.9994528079309752
```

```
In [62]: # get the analysis of roc auc curve for xgboost model
draw_roc_auc_curve(xgb_clf, "XGB Classifier", Ytest, Xtest)
```



```
In [63]: # confusion matrix analysis for xgboost classifier
draw_confusion_matrix(Ytest, xgb_clf.predict(Xtest), "XGB Classifier")
```

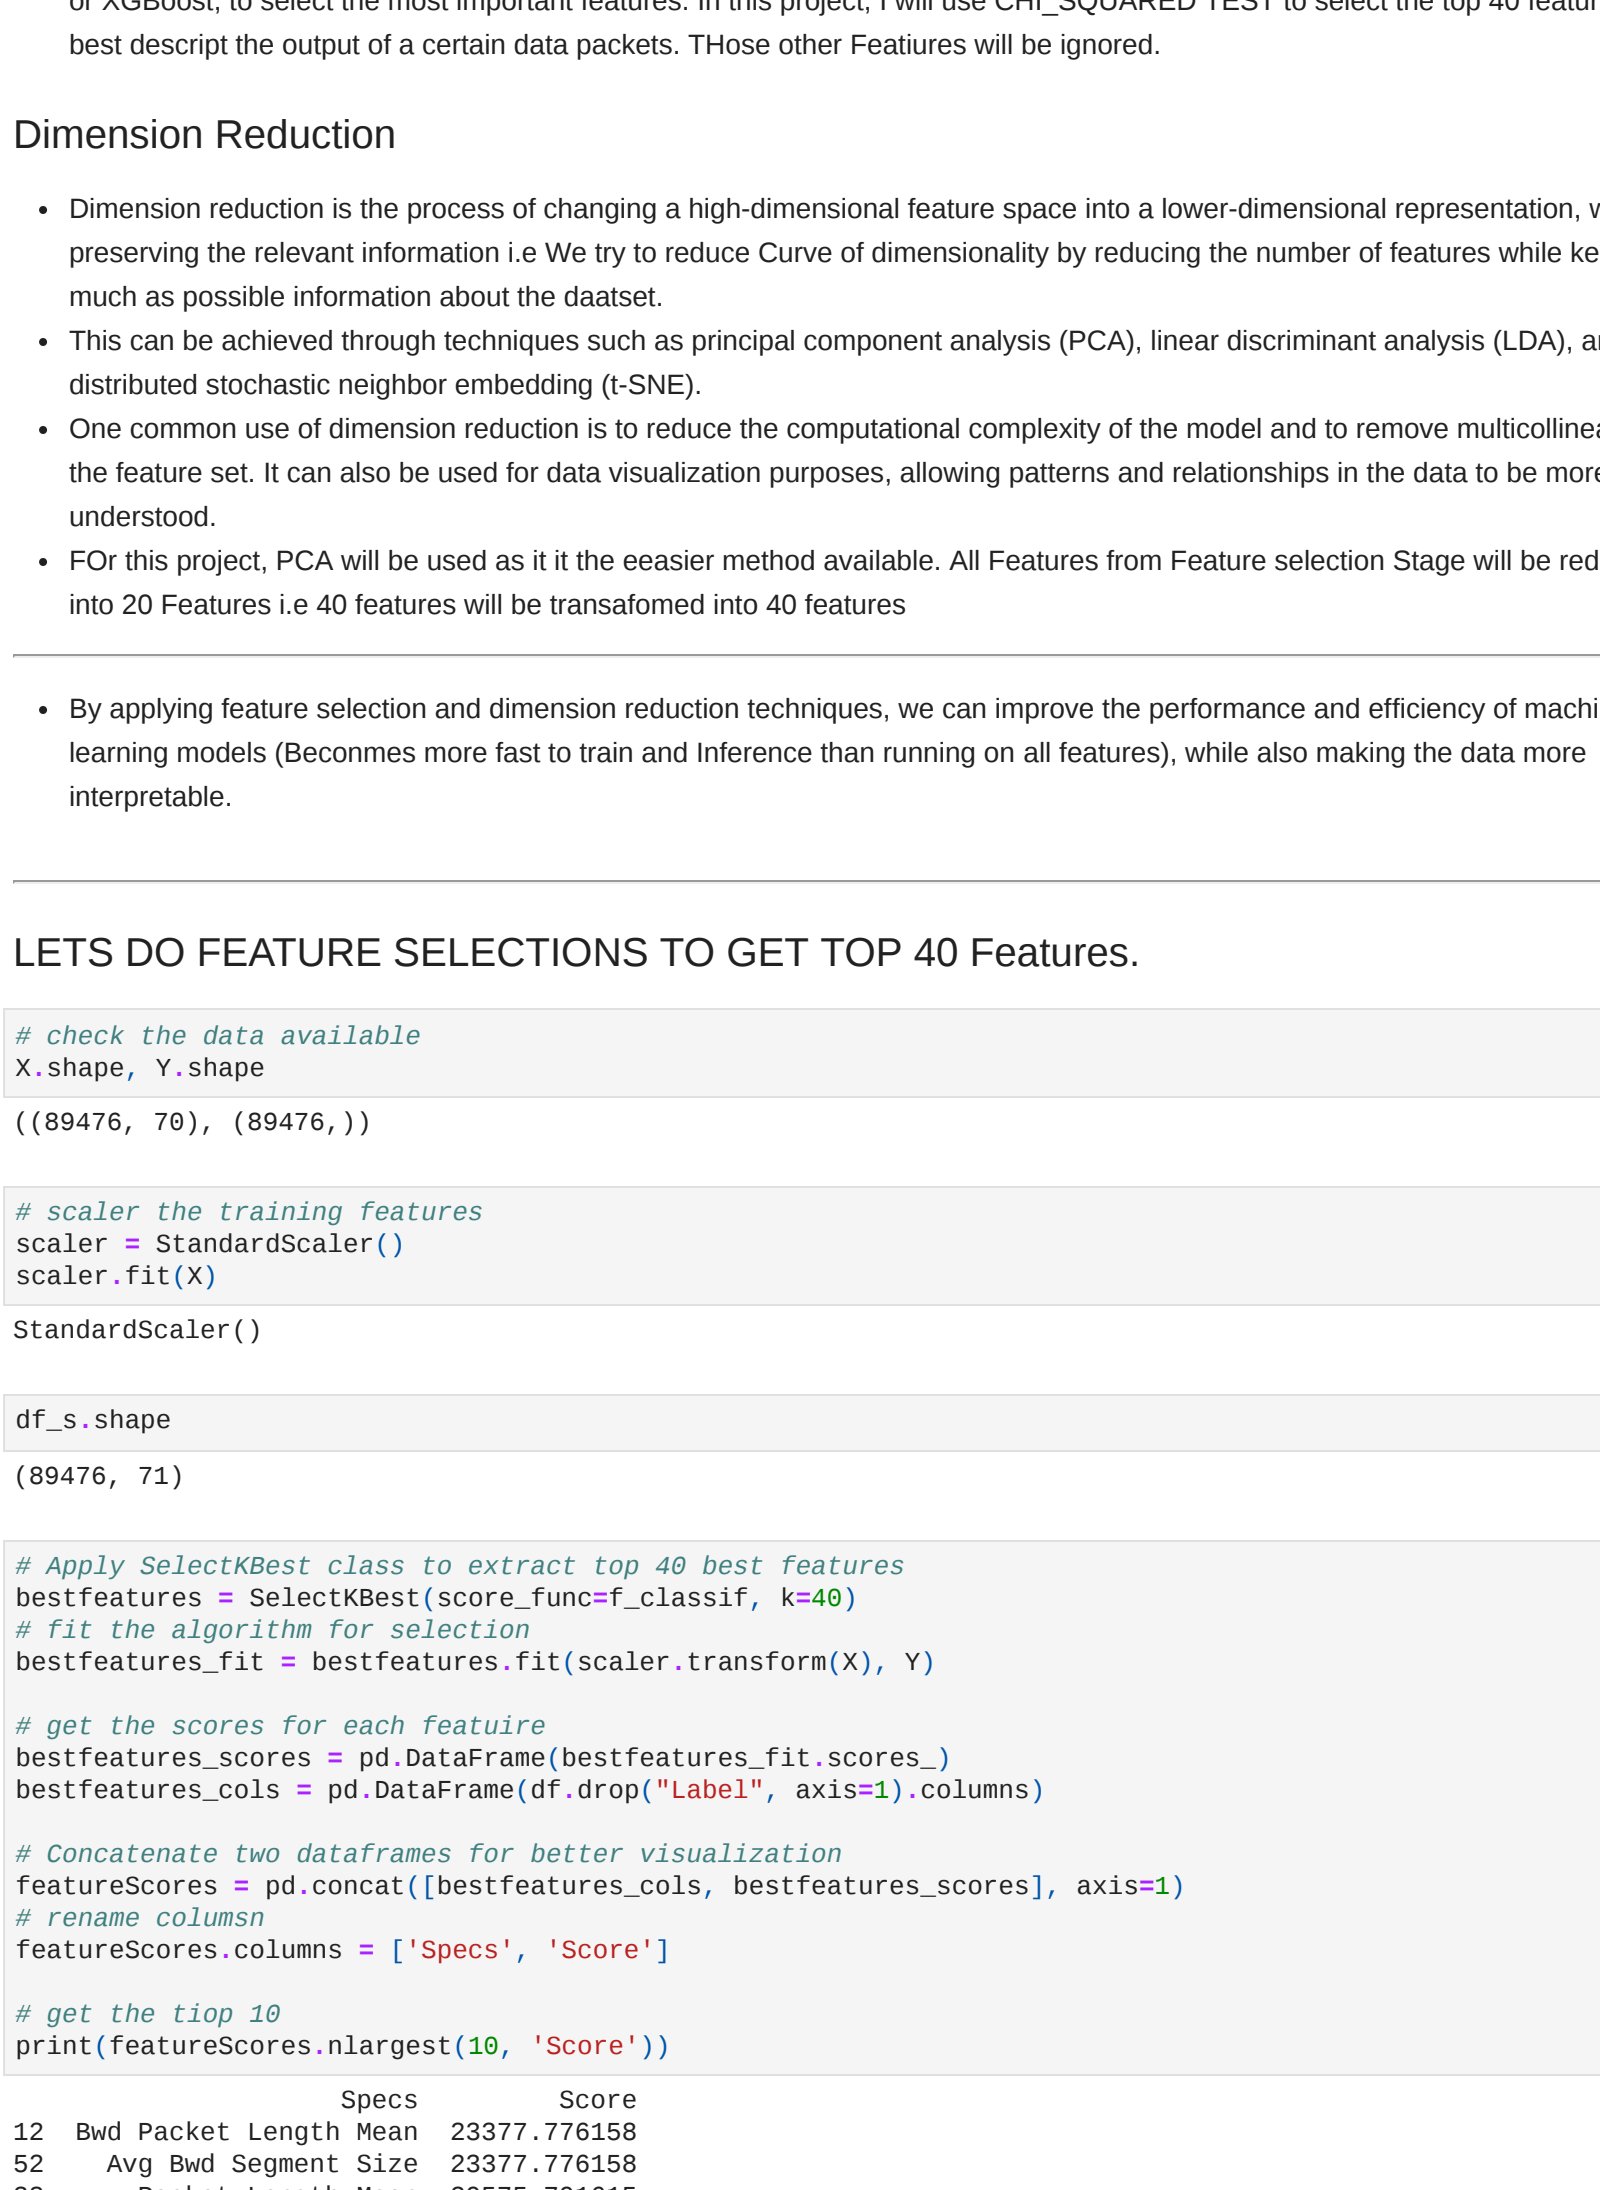
XGB Classifier Metrics Analysis Results

Accuracy is 0.9994523548740416
F1 Score is 0.9994528079309752
Recall score is 0.9994523548740416
Precision Score is 0.999451367653242

Outputting Classification Report Matrix Plot

	precision	recall	f1-score	support
BENIGN	1.00	1.00	1.00	4380
Bot	0.99	1.00	0.99	280
DDoS	1.00	1.00	1.00	2743
PortScan	1.00	1.00	1.00	3403
SSH_FTP_Patator	1.00	1.00	1.00	1976
accuracy		1.00	1.00	12782
macro avg	1.00	1.00	1.00	12782
weighted avg	1.00	1.00	1.00	12782

Confusion Plot for XGB Classifier's Performance



```
In [64]: # same random forest model
pd.to_pickle(rf_clf, "rf_clf.pkl")
```

Feature Selection and Dimension Reduction for CICIDS2017 DATASET.

- In Machine Learning, feature selection and dimension reduction are important techniques used to improve the performance and efficiency of a machine learning by reducing the number of input variables. The CICIDS2017 dataset provides a large number of features that can be used to train a machine learning model, however, not all of these features may be useful or necessary.

Feature Selection

- Feature selection is the process of selecting a subset of the most relevant features from the available feature set, based on a certain criterion. This helps to remove irrelevant and redundant features, and can improve the accuracy, speed, and interpretability of the model.
- One common method of feature selection is to use a univariate statistical test, such as the chi-squared test or ANOVA, to identify the most relevant features. Another method is to use feature importances generated by tree-based algorithms, such as Random Forest or XGBoost, to select the most important features. In this project, I will use CHI_SQUARED TEST to select the top 40 features that best describe the output of a certain data packets. Those other Features will be ignored.

Dimension Reduction

- Dimension reduction is the process of changing a high-dimensional feature space into a lower-dimensional representation, while preserving the relevant information. I use try to reduce Curve of dimensionality by reducing the number of features while keeping as much as possible information about the dataset.
- This can be achieved through techniques such as principal component analysis (PCA), linear discriminant analysis (LDA), and t-distributed stochastic neighbor embedding (t-SNE).
- One common use of dimension reduction is to reduce the computational complexity of the model and to remove multicollinearity in the feature set. PCA can also be used for data visualization purposes, allowing patterns and relationships in the data to be more easily understood.
- For this project, PCA will be used as it is the easier method available. All Features from Feature selection Stage will be reduced into 20 Features i.e 40 features will be transformed into 40 features

- By applying feature selection and dimension reduction techniques, we can improve the performance and efficiency of machine learning models. (Becomes more fast to train and inference than running on all features), while also making the data more interpretable.

LETS DO FEATURE SELECTIONS TO GET TOP 40 Features.

```
In [65]: # check the data available
X.shape, Y.shape
```

((89476, 70), (89476, 1))

```
In [66]: # scaler the training features
scaler = StandardScaler()
scaler.fit(X)
```

StandardScaler()

```
In [67]: df.s.shape
```

((89476, 71))

```
In [68]: # Apply SelectKBest class to extract top 40 best features
best_features = SelectKBest(score_func=f_classif, k=40)
# fit the algorithm for selection
best_features.fit = best_features.fit(scaler.transform(X), Y)
```

best_features_scores = pd.DataFrame(best_features.feature_scores_)

best_features_cols = pd.DataFrame(df.drop('label', axis=1).columns)

Concatenate two dataframes for better visualization

featuresScores = pd.concat([best_features_cols, best_features_scores], axis=1)

rename columns

featuresScores.columns = ['Specs', 'Score']

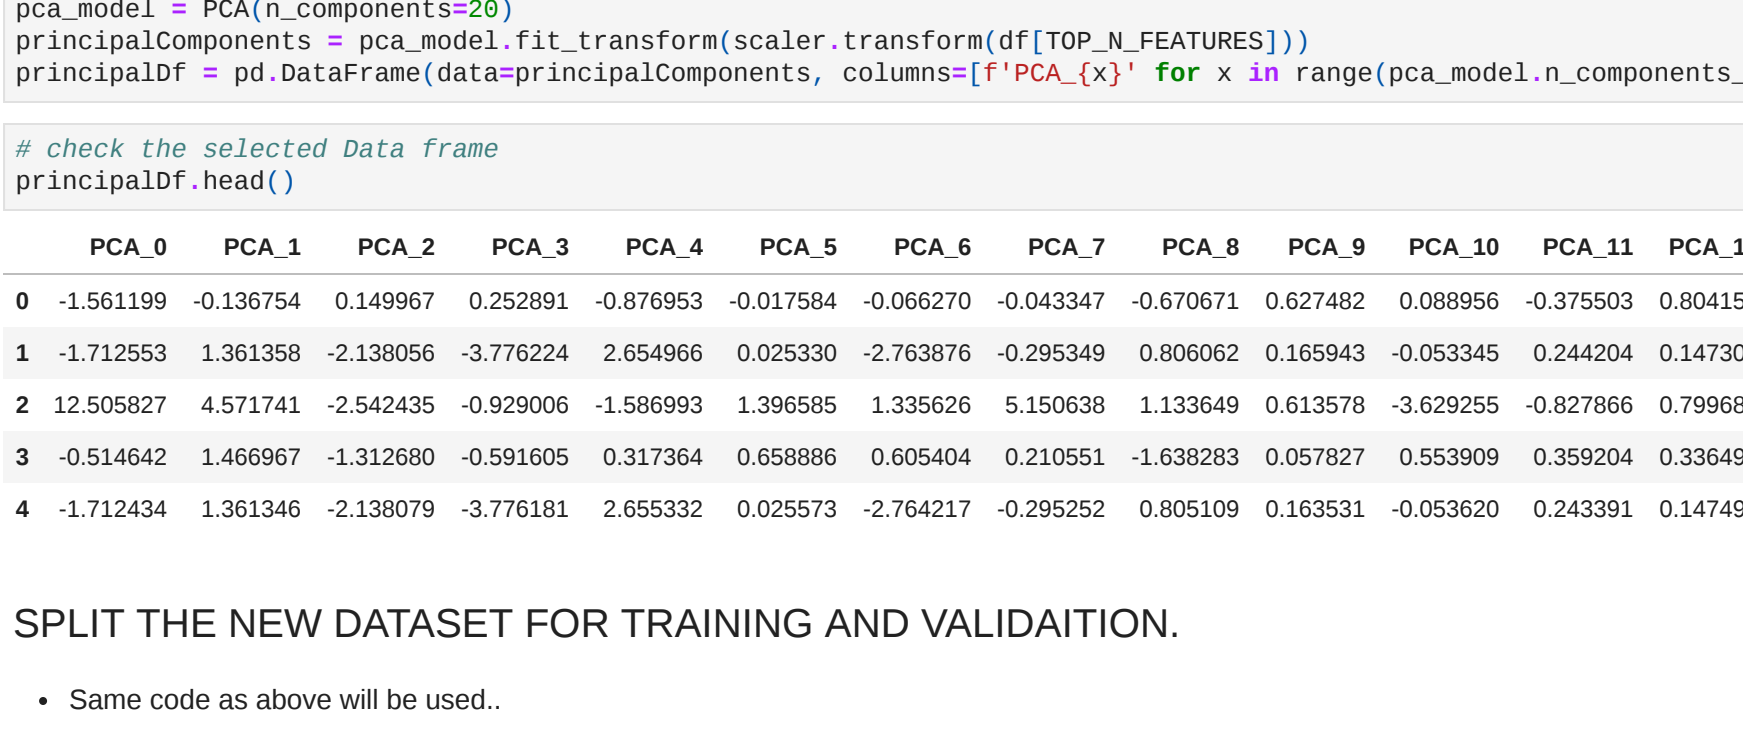
get the top 40

print(featuresScores.nlargest(40, 'Score'))

```
12  Bwd Packet Length Mean 23377.776158
52  Avg Bwd Segment Size 23377.776158
38  Packet Length Std 20575.791615
50  Average Packet Size 20561.120649
39  Packet Length Std 20407.351124
13  Bwd Packet Length Std 20072.057802
10  Bwd Packet Length Max 19959.085347
37  Max Packet Length 18237.738136
61  min_seg_size_forward 16060.970228
40  Packet Length Variance 14444.168514
```

```
In [ ]:
```

```
In [69]: # Get the indices sorted by most important to least important
indices = np.argsort(best_features.feature_scores_[::-1])
# To get your top 40 feature names
features = []
for i in range(40):
    features.append(df.drop('label', axis=1).columns[indices[i]])
# Now plot
plt.figure(figsize=(14,18))
plt.bar(features, best_features.feature_scores_[indices[range(40)]], color='r', align='center')
plt.xlabel('Features Name', fontsize=16, fontweight='bold')
plt.ylabel('Importance Score', fontsize=16, fontweight='bold')
plt.title('Features Importance for Top 40 Features', fontsize=19, fontweight='bold')
plt.show()
```



Results

- From the above, we have seen the top 40 features that contributes greatly to the type of packet.
- One of them is AVG BWD segment size which is the topmost features in determining what is the target type.
- It will only use the top 40 features to the next process of dimensionality reduction.
- The value 40 is selected randomly.

```
In [70]: TOP_N_FEATURES = features
print(TOP_N_FEATURES)
```

['Bwd Packet Length Mean', 'Avg Bwd Segment Size', 'Packet Length Mean', 'Average Packet Size', 'Packet Length Std', 'Bwd Packet Length Std', 'Bwd Packet Length Max', 'Max Packet Length', 'min_seg_size_forward', 'Packet Length Variance', 'PSH Flag Count', 'Min Packet Length', 'Bwd Packet Length Min', 'ACK Flag Count', 'SYN Flag Count', 'Fwd PSH Flags', 'Flow IAT Std', 'Init_Win_bytes_forward', 'Flow IAT Max', 'Idle Max', 'Flow IAT Max', 'URG Flag Count', 'Destination Port', 'Idle Mean', 'Bwd Packets/s', 'Fwd Packets/s', 'Flow Du nt', 'Fwd IAT Total', 'Idle Std', 'Idle Min', 'Flow IAT Mean', 'Bwd IAT Total', 'DownUp Ratio', 'Bwd Header Length', 'Fwd Packet Length Mean', 'Avg Bwd Segment Size', 'Fwd Header Length', 'Fwd Header Length.1', 'Bwd IAT Max']

```
In [71]: len(TOP_N_FEATURES)
```

40

PCA FOR DIMENSIONS REDUCTION.

- Below is implementation of PCA that takes the whole 40 features data selected above and reduce it into an array of 20 features (just a random selection).
- Here is the code

```
In [72]: # define scaler to transform the dataset
scaler = StandardScaler()
scaler.fit(df[TOP_N_FEATURES])

pca_model = PCA(n_components=20)
principalComponents = pca_model.fit_transform(scaler.transform(df[TOP_N_FEATURES]))
principalDf = pd.DataFrame(data=principalComponents, columns=['PCA_x'], for x in range(pca_model.n_components))
```

```
In [73]: # check the selected data frame
print(principalDf.head())
```

```
Out[73]:
```

```
PCA.0 PCA.1 PCA.2 PCA.3 PCA.4 PCA.5 PCA.6 PCA.7 PCA.8 PCA.9 PCA.10 PCA.11 PCA.12
0 -1.561199 -0.130154 0.149967 0.252891 -0.769953 -0.017584 -0.066270 -0.043347 -0.670671 0.627482 0.088956 -0.375503 0.004151
1 -1.712953 1.361358 -2.138056 -3.776224 2.654966 0.025330 -2.763876 -0.206347 0.806062 0.165943 -0.053345 0.244204 0.147409
2 12.509827 4.571741 -2.542426 0.072006 1.589391 1.396285 1.326526 5.150638 1.133649 0.613578 -3.629255 0.827865 0.799689
3 -0.514642 1.466867 -1.322880 -0.591005 0.317964 0.669886 0.605404 0.210951 -1.638293 0.057827 0.555909 0.395924 0.336485
4 -1.712434 1.361346 -2.138079 -3.776181 2.655332 0.025373 -2.764217 -0.205252 0.805109 0.165531 -0.053620 0.242391 0.147495
```

SPLIT THE NEW DATASET FOR TRAINING AND VALIDATION.

- Same code as above will be used.

```
In [74]: # we copied same code from above but these time we use training features that is transformed
for train_index, test_index in skf.split(principalDf, Y):
    X_train, X_test = X_train.index, X_test.index
    Y_train, Y_test = Y_train.index, Y_test.index
    reshaped_y_train = np.asarray(Y_train).reshape(-1, 1)
    reshaped_y_test = np.asarray(Y_test).reshape(-1, 1)
print('X_train length: ', len(X_train)) # To check if splits worked
print('Y_train length: ', len(Y_train))
print('X_test length: ', len(X_test))
print('Y_test length: ', len(Y_test))

X_train length: 76694
Y_train length: 76694
X_test length: 12782
Y_test length: 12782
```

Retraining Models to observe their Performace.

- I will use similar models as above Except that I will only Test Logistic Regression and Random Forest and Later Build An Ensemble with the four models and check with the same metrics.

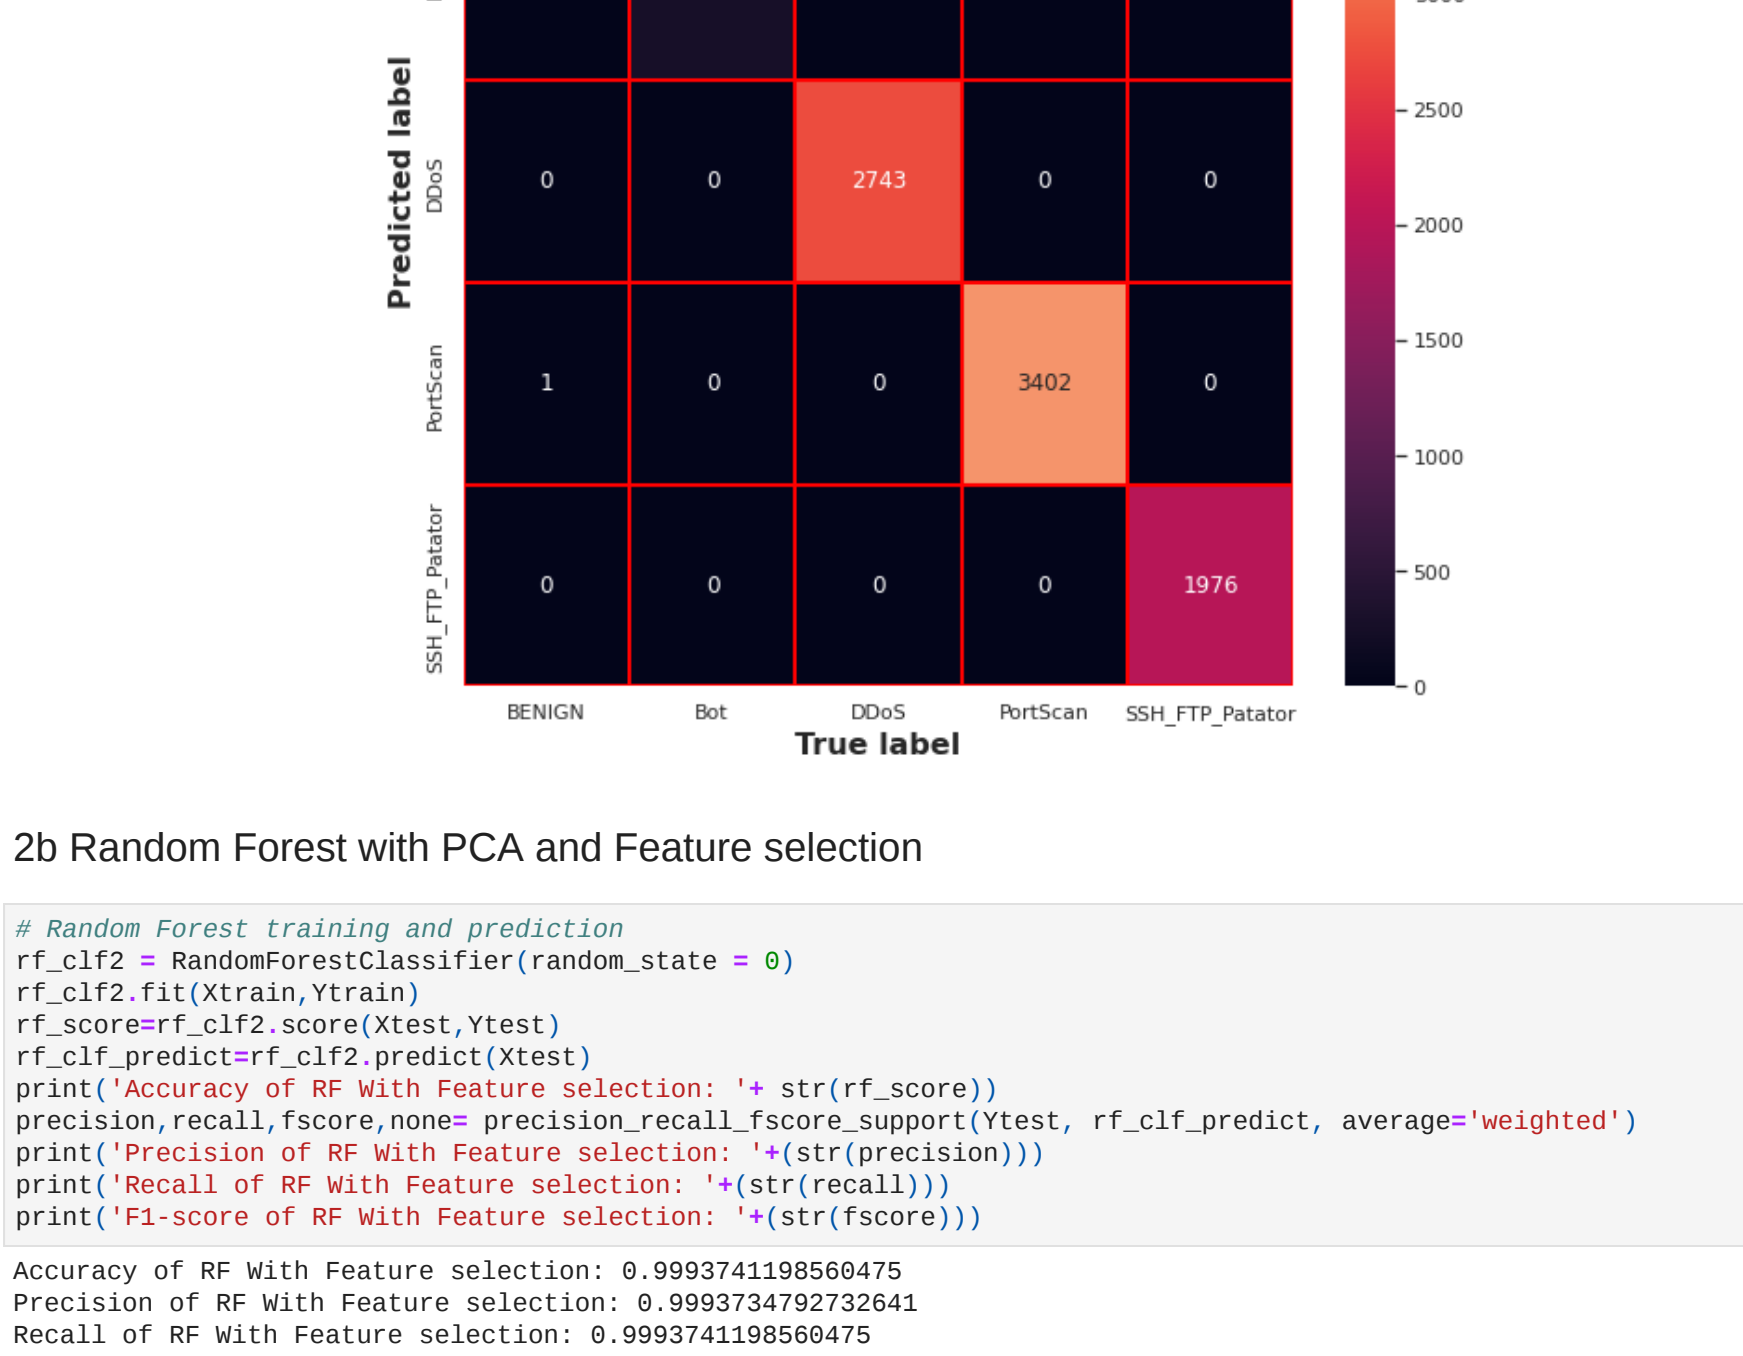
1b. Logistic regression with Feature selection and Dimensionality reduction

```
In [75]: # Logistic Regression Model
log_clf2 = LogisticRegressionClassifier(random_state = 0)
# train the model
log_clf2.fit(X_train,Y_train)

# get model score and predictions
log_score=log_clf2.score(Xtest,Ytest)
log_preds=log_clf2.predict(Xtest)
print('Accuracy of Logistic Regression Classifier with Feature selection: %s'%(str(log_score)))
precision,recall,f1score,none: precision_recall_fscore_support(Ytest, log_preds, average='weighted')
print('Precision of Logistic Regression Classifier with Feature selection: %s'%(str(precision)))
print('Recall of RF With Feature selection: %s'%(str(recall)))
print('F1-score of RF With Feature selection: %s'%(str(f1score)))

Accuracy of Logistic Regression Classifier with Feature selection: 0.9989047097480832
Precision of Logistic Regression Classifier with Feature selection: 0.9989024261350944
Recall of Logistic Regression Classifier with Feature selection: 0.9989047097480832
F1-score of Logistic Regression Classifier with Feature selection: 0.9989037492373268
```

```
In [76]: # Feature reduction result
# get the analysis of roc auc curve for logistic regression model
draw_roc_auc_curve(log_clf2, "Logistic Regression Classifier with Feature selection", Ytest, Xtest)
```



```
In [77]: # confusion matrix analysis for logistic regression model
draw_confusion_matrix(Ytest, log_clf2.predict(Xtest), "Logistic Regression Classifier with Feature selection")
```

Logistic Regression Classifier Metrics Analysis Results

Accuracy is 0.9989047097480832
F1 Score is 0.9989017492373268
Recall score is 0.9989047097480832
Precision Score is 0.9989024261350944

Outputting Classification Report Matrix Plot

	precision	recall	f1-score	support
BENIGN	1.00	1.00	1.00	4380
Bot	0.99	0.97	0.98	280
DDoS	1.00	1.00	1.00	2743
PortScan	1.00	1.00	1.00	3403
SSH_FTP_Patator	1.00	1.00	1.00	1976
accuracy		1.00	1.00	12782
macro avg	1.00	0.99	1.00	12782
weighted avg	1.00	1.00	1.00	12782

Confusion Plot for Logistic Regression Classifier with Feature selection's Performance

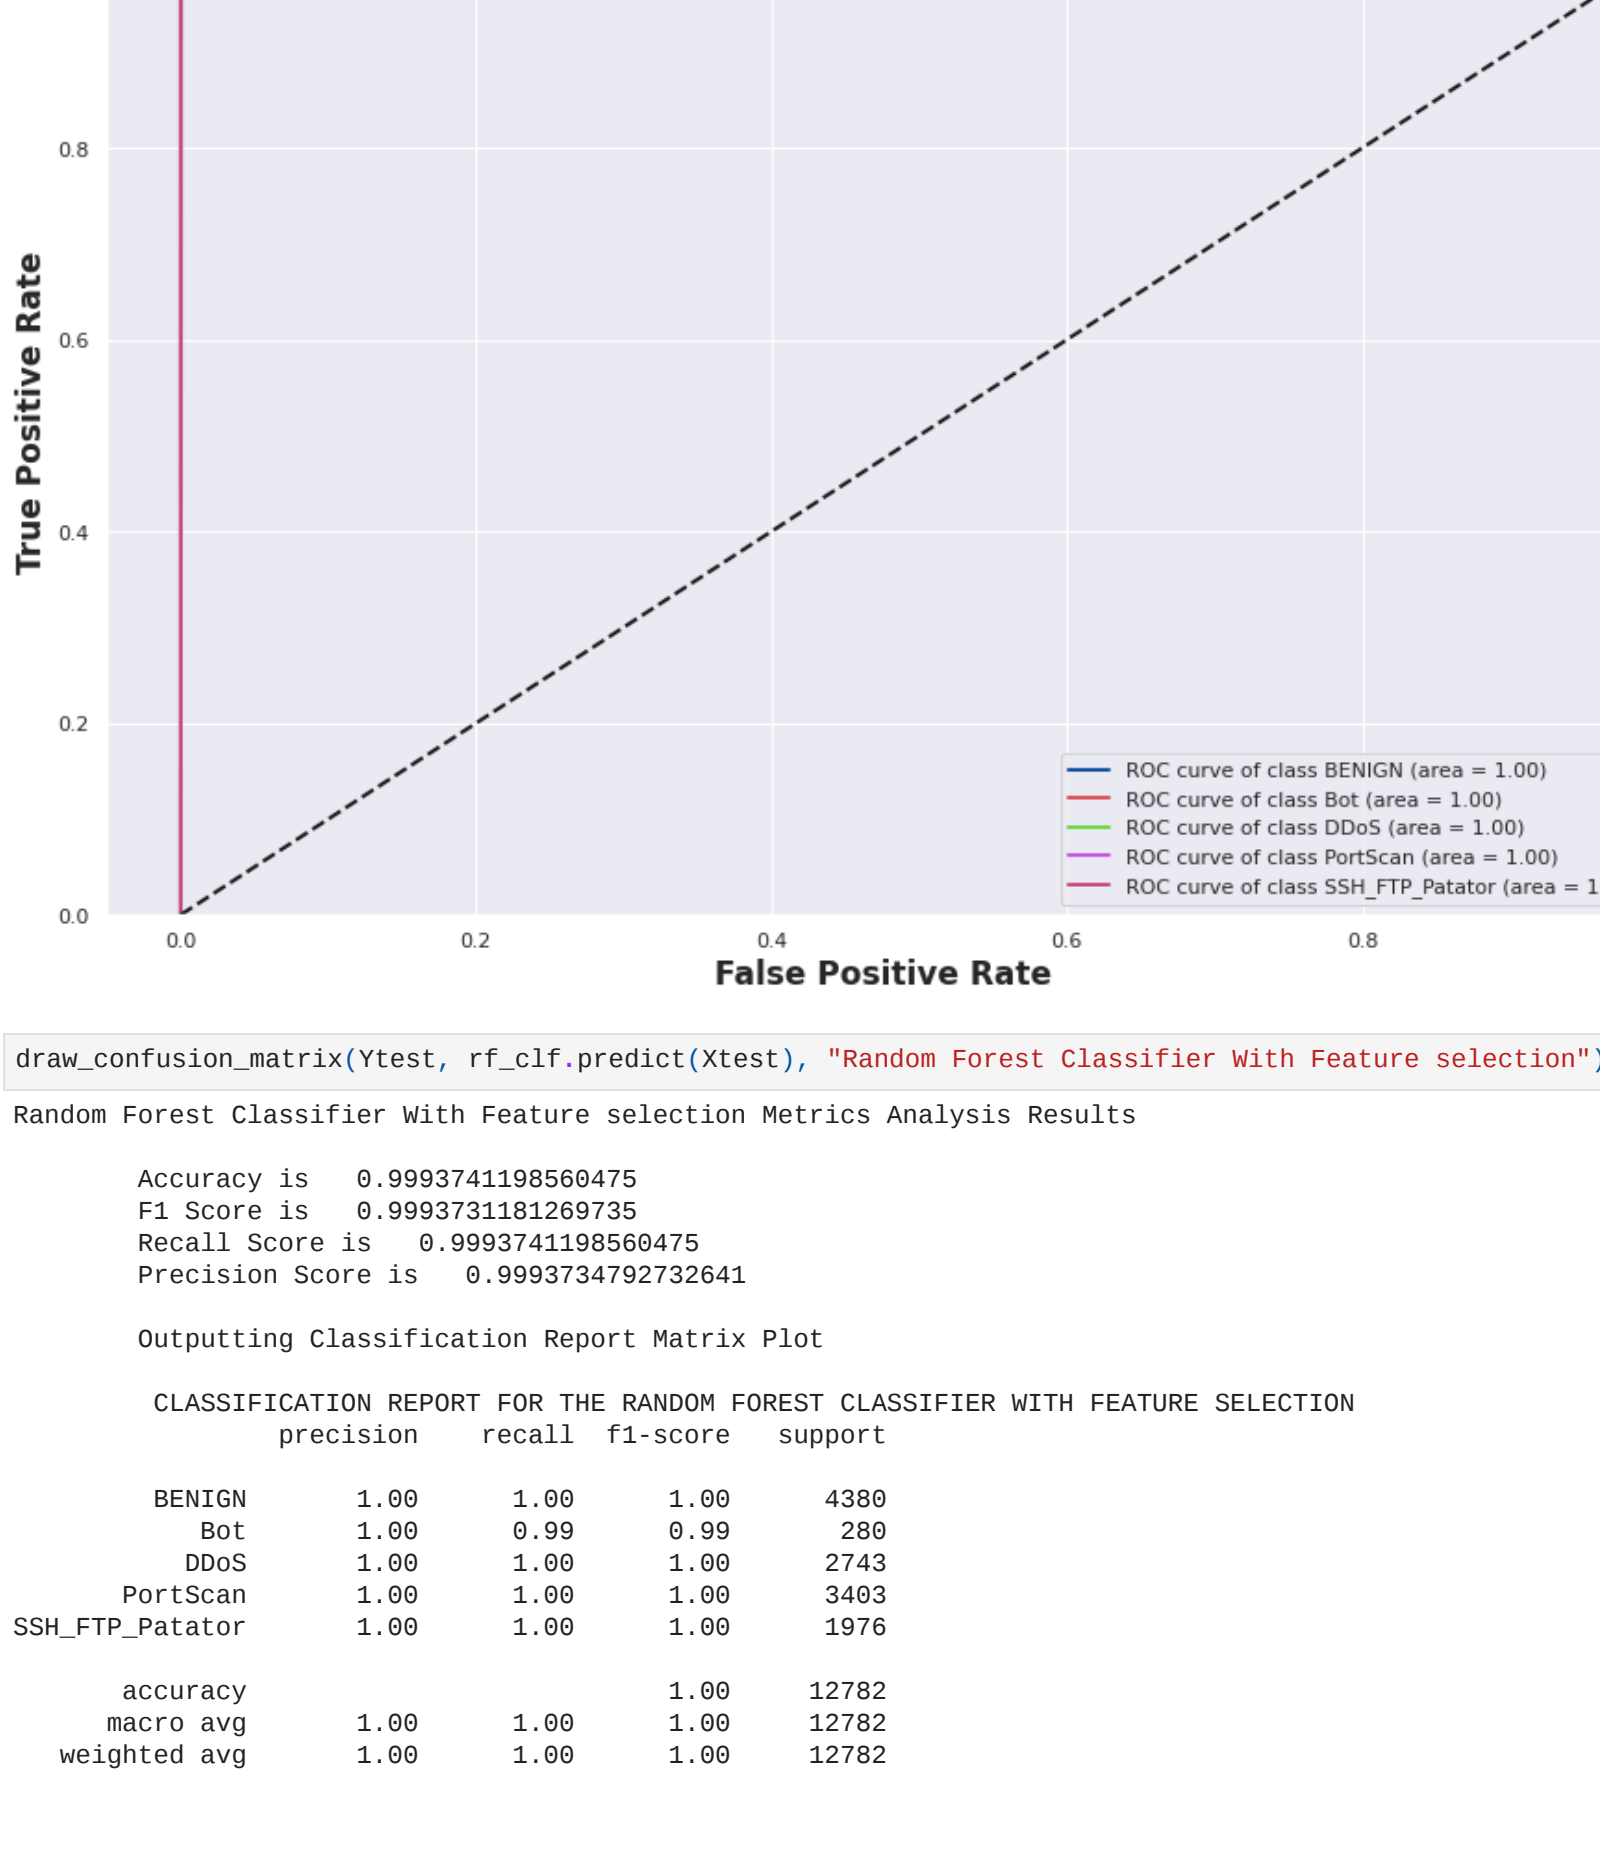
2b Random Forest with PCA and Feature selection

```
In [78]: # Random Forest training and prediction
rf_clf2 = RandomForestClassifier(random_state = 0)
rf_clf2.fit(X_train,Y_train)
rf_score=rf_clf2.score(Xtest,Ytest)
rf_clf_predict=rf_clf2.predict(Xtest)
xgb_predict=xgb_clf.predict(Xtest)
print('Accuracy of RF: %s'%(str(rf_score)))
precision,recall,f1score,none: precision_recall_fscore_support(Ytest, rf_clf_predict, average='weighted')
print('Precision of RF With Feature selection: %s'%(str(precision)))
print('Recall of RF With Feature selection: %s'%(str(recall)))
print('F1-score of RF With Feature selection: %s'%(str(f1score)))

Accuracy of RF With Feature selection: 0.9993741198560475
Precision of RF With Feature selection: 0.9993734792732641
Recall of RF With Feature selection: 0.9993741198560475
F1-score of RF With Feature selection: 0.9993731181269735
```

```
In [79]: # RF WITH FEATURE SELECTION
# get the analysis of roc auc curve for random forest
draw_roc_auc_curve(rf_clf2, "Random Forest Classifier With Feature selection", Ytest, Xtest)
```


Receiver operating characteristic for Random Forest Classifies With Feature selection



```
In [80]: draw_confusion_matrix(Ytest, rf_clf.predict(Xtest), "Random Forest Classifier With Feature selection")

Random Forest Classifier With Feature selection Metrics Analysis Results

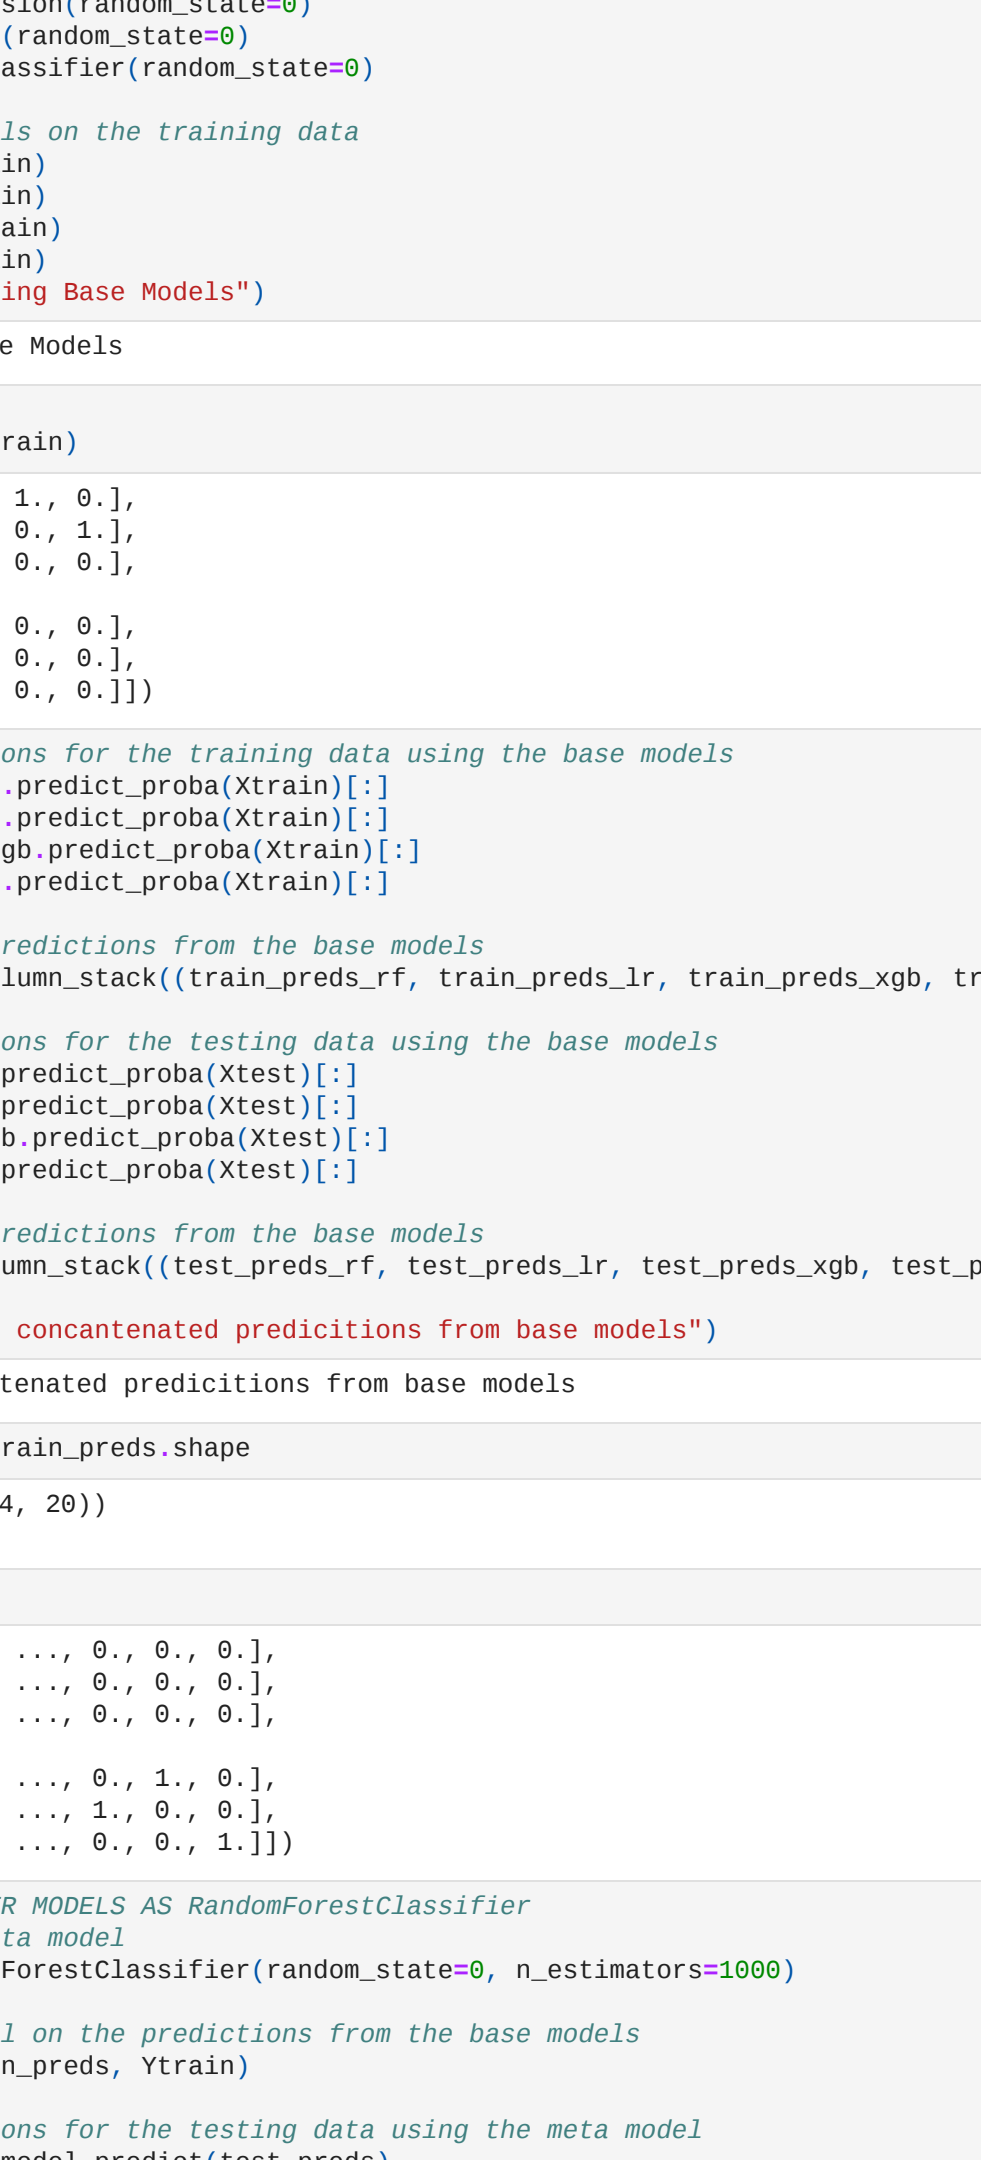
Accuracy is 0.9993741198560475
F1 Score is 0.9993731161269735
Recall Score is 0.9993741198560475
Precision Score is 0.9993734792732641

Outputting Classification Report Matrix Plot

CLASSIFICATION REPORT FOR THE RANDOM FOREST CLASSIFIER WITH FEATURE SELECTION
precision recall f1-score support
BENIGN 1.00 1.00 1.00 4380
Bot 1.00 0.99 0.99 280
DDoS 1.00 1.00 1.00 2743
PortScan 1.00 1.00 1.00 3403
SSH_FTP_Patator 1.00 1.00 1.00 1976

accuracy 1.00 1.00 1.00 12782
macro avg 1.00 1.00 1.00 12782
weighted avg 1.00 1.00 1.00 12782
```

Confusion Plot for Random Forest Classifier With Feature selection's Performance



FINALLY.

ENSEMBLING THROUGH STACKING.

- A stacking model is an ensemble learning technique where multiple models are combined to make a final prediction. In this case, I will stack random forest, logistic regression, XGBoost, and decision trees to build a more powerful model.
- The four base models (random forest, logistic regression, XGBoost, and decision trees) will be trained on the training data, and their predictions are generated for both the training and testing data. These predictions are then concatenated and used as input for the meta model (logistic regression in this case), which is also trained on the training data. Finally, the accuracy of the stacking model is calculated using the testing data. Below is the implementation

```
In [81]: # Initialize the base models
rf = RandomForestClassifier(n_estimators=100, random_state=0)
lr = LogisticRegression(random_state=0)
xgb = XGBClassifier(random_state=0)
dt = DecisionTreeClassifier(random_state=0)

# Fit the base models on the training data
rf.fit(Xtrain, Ytrain)
lr.fit(Xtrain, Ytrain)
xgb.fit(Xtrain, Ytrain)
dt.fit(Xtrain, Ytrain)
print("Done Predicting Base Models")

Done Predicting Base Models
```

```
In [86]: # get probs
rf.predict_proba(Xtrain)

Out[86]: array([[0., 0., 0., 1., 0.],
        [0., 0., 0., 0., 1.],
        [0., 0., 1., 0., 0.],
        ...,
        [0., 0., 1., 0., 0.],
        [0., 0., 1., 0., 0.],
        [0., 0., 1., 0., 0.]])
```

```
In [87]: # Generate predictions for the training data using the base models
train_preds_rf = rf.predict_proba(Xtrain)[:,1]
train_preds_lr = lr.predict_proba(Xtrain)[:,1]
train_preds_xgb = xgb.predict_proba(Xtrain)[:,1]
train_preds_dt = dt.predict_proba(Xtrain)[:,1]

# Concatenate the predictions from the base models
train_preds = np.column_stack((train_preds_rf, train_preds_lr, train_preds_xgb, train_preds_dt))

# Generate predictions for the testing data using the base models
test_preds_rf = rf.predict_proba(Xtest)[:,1]
test_preds_lr = lr.predict_proba(Xtest)[:,1]
test_preds_xgb = xgb.predict_proba(Xtest)[:,1]
test_preds_dt = dt.predict_proba(Xtest)[:,1]

# Concatenate the predictions from the base models
test_preds = np.column_stack((test_preds_rf, test_preds_lr, test_preds_xgb, test_preds_dt))

print("Done getting concatenated predictions from base models")

Done getting concatenated predictions from base models
```

```
In [89]: test_preds.shape, train_preds.shape

Out[89]: ((12782, 20), (76694, 20))

In [90]: test_preds

Out[90]: array([[0., 1., 0., ..., 0., 0., 0.],
        [0., 1., 0., ..., 0., 0., 0.],
        [0., 1., ..., 0., 0., 0.],
        ...,
        [0., 0., 0., ..., 0., 1., 0.],
        [0., 0., 1., ..., 1., 0., 0.],
        [0., 0., 1., ..., 0., 1.]])
```

```
In [91]: # BUILDING A STACKER MODELS AS RandomForestClassifier
# Initialize the meta model
meta_model = RandomForestClassifier(random_state=0, n_estimators=1000)

# Fit the meta model on the predictions from the base models
meta_model.fit(train_preds, Ytrain)

# Generate predictions for the testing data using the meta model
Final_preds = meta_model.predict(test_preds)
```

```
In [95]: # EVALUATE THE META MODEL TO SEE ITS PERFORMANCE

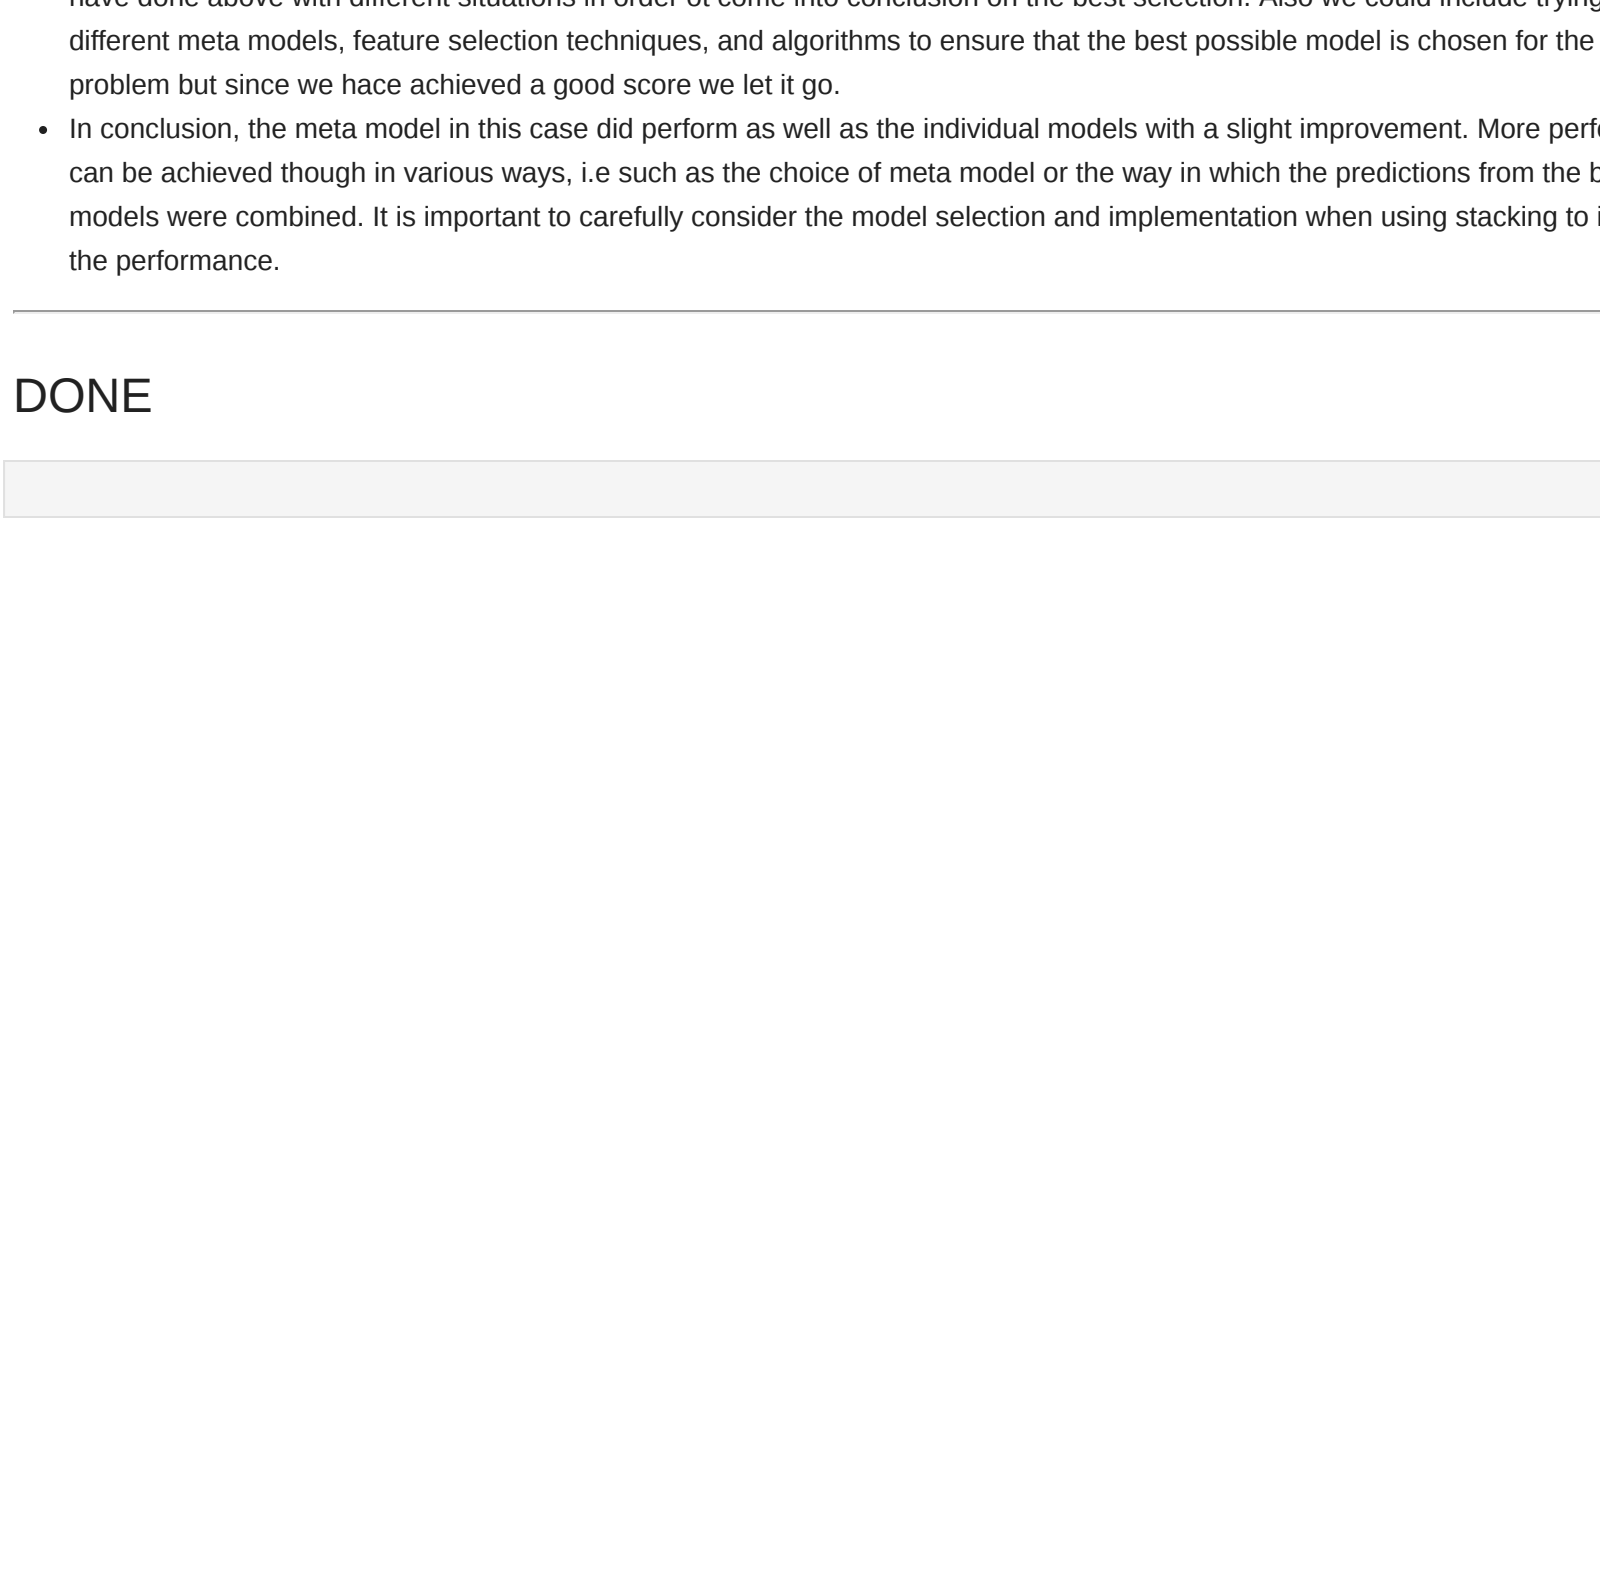
# Calculate the accuracy of the
accuracy = accuracy_score(Ytest, final_preds)
print("Accuracy of Stacking Model: {:.2f}%".format(accuracy*100))

Accuracy of Stacking Model: 99.98%
```

EVALUATING STACKED MODELS

```
In [93]: # STACKED MODEL EVALUATION WITH FEATURE SELECTION
# FPR analysis for meta model
draw_roc_auc_curve(meta_model, "Random Forest Classifies AS META MODEL", Ytest, test_preds)
```

Receiver operating characteristic for Random Forest Classifies AS META MODEL



```
In [94]: # plot confusion matrix for meta model
draw_confusion_matrix(Ytest, final_preds, "Random Forest Classifies AS META MODEL")

Random Forest Classifies AS META MODEL Metrics Analysis Results

Accuracy is 0.999688249109298
F1 Score is 0.9996885738726254
Recall Score is 0.999688249109298
Precision Score is 0.999688498909999

Outputting Classification Report Matrix Plot

CLASSIFICATION REPORT FOR THE RANDOM FOREST CLASSIFIES AS META MODEL
precision recall f1-score support
BENIGN 1.00 1.00 1.00 4380
Bot 1.00 0.99 1.00 280
DDoS 1.00 1.00 1.00 2743
PortScan 1.00 1.00 1.00 3403
SSH_FTP_Patator 1.00 1.00 1.00 1976

accuracy 1.00 1.00 1.00 12782
macro avg 1.00 1.00 1.00 12782
weighted avg 1.00 1.00 1.00 12782
```

Confusion Plot for Random Forest Classifies AS META MODEL's Performance



CONCLUSION.

- The results of a stacking model using Random Forest as the meta model show that the individual models (Random Forest, Logistic Regression, XGBoost, Decision tree classifier) were outperformed better by the stacking model. The accuracy, F1 score, recall score, and precision score of the individual models were abit lower compared to the stacking model. The accuracy of the Random Forest model was 0.999, the Logistic Regression model was 0.998, the Decision tree model was 0.999, and the XGBoost model was 0.999. However, the stacking model had an accuracy of only 1.000. Where only 3 classes were not predicted accurately.
- Based on the results of the individual models and the stacking model, it is clear that the Random Forest, Logistic Regression, and XGBoost and Meta Model models performed better. To determine which model to use, it would be necessary to consider additional factors such as the business problem and the specific requirements for the model.
- Since the goal was to achieve the best model, it seems that using the Meta Model of Random Forest or XGBoost model would be the best choice.
- Normally it is important to thoroughly evaluate and compare the performance of multiple models to make an informed decision as we have done above with different situations in order to come into conclusion on the best selection. Also we could include trying different meta models, feature selection techniques, and algorithms to ensure that the best possible model is chosen for the given problem but since we have achieved a good score we let it go.
- In conclusion, the meta model in this case did perform as well as the individual models with a slight improvement. More performance can be achieved though in various ways, i.e. such as the choice of meta model or the way in which the predictions from the base models were combined. It is important to carefully consider the model selection and implementation when using stacking to improve the performance.

DONE

```
In [ ]:
```